
IVY-FAKE: A Unified Explainable Framework and Benchmark for Image and Video AIGC Detection

Supplementary Material

Anonymous Author(s)

Affiliation

Address

email

1 In this section, we present additional experiments to thoroughly evaluate the proposed approach,
2 along with implementation details for the distillation prompts. The experiments specifically focus
3 on the reasoning capabilities of the AGIC detection task. Rather than merely providing final real or
4 fake predictions, the methods here also generate detailed explanations describing the rationale behind
5 each prediction.

6 The evaluation is primarily conducted on the proposed explainability benchmark, the Ivy-Fake dataset.
7 Additionally, we assess the generalization capability of our method on a broader video understanding
8 dataset.

9 A Performance Evaluation of Explainability Datasets

10 We conduct experiments to evaluate the explanation capabilities of various models. Specifically,
11 we assess the similarity between model-generated reasoning and reference annotations using the
12 ROUGE-L score (Lin, 2004), which measures the longest common subsequence and reflects token-
13 level overlap. Additionally, we adopt the LLM-as-a-judge paradigm (Zheng et al., 2023), evaluating
14 responses across four dimensions: Completeness, Relevance, Level of Detail, and Explanation
15 Quality. Each response is scored using GPT-4o mini (Achiam et al., 2023) under a unified prompt
16 that instructs the model to act as an impartial judge, assigning scores from 1 to 5. To ensure reliability,
17 each response is rated across five independent rounds, with the final score computed as the average.

Model	Image							Video						
	Auto Metrics				GPT Assisted			Auto Metrics				GPT Assisted		
	Acc	F1	ROUGE-L	SIM	Com./Rel./Det./Exp.	AVG		Acc	F1	ROUGE-L	SIM	Com./Rel./Det./Exp.	AVG	
<i>Closed-source MLLMs</i>														
GPT-4o	0.766	0.759	0.155	0.683	3.69/3.79/3.67/3.80	3.74		0.828	0.813	0.150	0.682	4.02/4.13/3.99/4.12	4.07	
Gemini 2.5 flash	0.668	0.656	0.223	0.709	3.29/3.33/3.27/3.33	3.30		0.787	0.784	0.188	0.678	3.85/3.93/3.82/3.92	3.88	
GPT4o-mini	0.653	0.650	0.149	0.645	3.21/3.26/3.21/3.26	3.23		0.685	0.672	0.134	0.645	3.38/3.42/3.37/3.42	3.40	
<i>Open-source MLLMs</i>														
InternVL3-8B	0.614	0.605	0.159	0.680	3.04/3.07/3.04/3.07	3.05		0.632	0.616	0.165	0.629	3.13/3.16/3.12/3.16	3.14	
Qwen2.5-VL-7B	0.556	0.516	0.199	0.660	2.73/2.77/2.73/2.77	2.75		0.589	0.527	0.207	0.621	2.89/2.94/2.89/2.89	2.91	
Phi-3.5-Vision	0.560	0.555	0.092	0.366	2.66/2.72/2.66/2.72	2.69		0.559	0.479	0.001	0.046	2.78/2.79/2.78/2.79	2.79	
Ivy-xDet	0.901	0.894	0.213	0.710	4.39/4.21/4.33/4.54	4.40		0.945	0.945	0.303	0.776	4.50/4.71/4.42/4.68	4.58	

Table 1: Performance comparison of models on image and video tasks. “Auto Metrics” include Acc, F1, ROUGE-L, and SIM scores. “GPT Assisted” includes four subjective criteria: Comprehensive-ness, Relevance, Detail, and Explanation, as well as their average.

18 As shown in Table 1, we compare our model with several leading multimodal large language
19 models (LLMs), including Qwen2.5-7B (Bai et al., 2025), InternVL2.5-8B (Chen et al., 2024a,b),

20 Phi-3.5-Vision, GPT-4V (Achiam et al., 2023), and Gemini 2.5 (Team et al., 2023). The results
 21 demonstrate that our approach not only achieves higher accuracy but also provides more transparent
 22 and informative explanations than all baseline models.

23 B Video Understanding Models and Evaluation on General Benchmarks

24 We further evaluate the proposed model on several video understanding benchmarks to assess its
 25 generalization capability. Specifically, we compare it against five lightweight, general-purpose video
 26 understanding models, which are VideoLLaMA3, Qwen2-VL 2B, Qwen2.5-VL-3B, InternVL2.5-2B,
 27 and InternVL3-2B, across four benchmarks, i.e., MLVU (dev), PerceptionTest, LongVideo, and
 28 VideoMME.

29 As shown in Table 2, our model, Ivy-Video-3B, consistently outperforms all competing methods
 30 across these benchmarks. These results highlight the strong generalization ability of our model, which,
 31 although designed for AGIC detection, achieves high accuracy across a diverse set of general-purpose
 32 video understanding tasks.

Task	VideoLLaMA3	Qwen2-VL 2B	Qwen2.5-VL-3B	InternVL2.5-2B	InternVL3-2B	Ivy-Video-3B
MLVU (dev)	65.4	62.7	68.2	58.9	64.2	68.87
PerceptionTest	68	53.9	66.9	66.3	-	72.73
LongVideo	54.3	44.7	50.2	48.7	51.2	55.25
VideoMME	59.6	55.6	61.5	51.9	58.9	62.3

Table 2: Transposed performance table of various video understanding models on general video QA benchmarks.

33 C Effect of Incorporating Human-Annotated Labels via gemini 2.5 pro on Accuracy

34
 35 To assess the impact of human-annotated labels on model performance, we compare the accuracy
 36 of final conclusion predictions under two settings: (i) with labels incorporated via the gemini 2.5
 37 pro, and (ii) without labels. The evaluation was conducted on about 1,000 examples from the test set.

Annotation Setting	Accuracy (Acc)
With Label	1.000
Without Label	0.785

Table 3: Accuracy of conclusion prediction with and without incorporating labels.

38 As shown in Table 3, incorporating ground-truth labels results in a substantial performance gain,
 39 yielding perfect accuracy (1.000), compared to 0.785 without labels. The drastic performance gap
 40 suggests potential limitations in label-free or weakly supervised setups when applied to tasks requiring
 41 fine-grained semantic understanding.

42 D Prompts

43 Here we provide the prompts that are mainly used in this study. As illustrated by the following
 44 figures, there are five distillation prompts distillation we used in this paper that mainly can be divided
 45 into the following three parts:

46
 47 **1.Prompt Template for Image Data Distillation:** Since image data consists of a single frame, it
 48 can be treated as a static instance. Therefore, AIGC detection mainly focuses on identifying spatial
 49 anomalies. Detail prompt can be found in Figure 1 and 3.

50 **2.Prompt Template for Video Data Distillation:** Compared to images, video inputs provide
 51 continuous multi-frame context. This allows for detection along both spatial and temporal anomaly
 52 dimensions. etail prompt can be found in 2 and Figure 4.

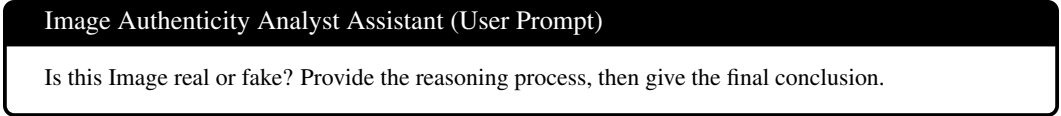


Figure 1: User Prompt Template for Image Data Distillation

53 **3.GPT Assisted Evaluation Prompt:** To assess the quality of model outputs, we design a GPT-based
54 evaluator prompt that scores responses across four dimensions: Completeness, Relevance, Level of
55 Detail, and Explanation. The evaluator receives a structured pair of GroundTruth and ModelOutput,
56 each containing a <think> section (reasoning) and a <conclusion> (final judgment). The model must
57 return a structured JSON object with integer scores (1–5) for each dimension. The prompt is provided
58 in Figure 5.

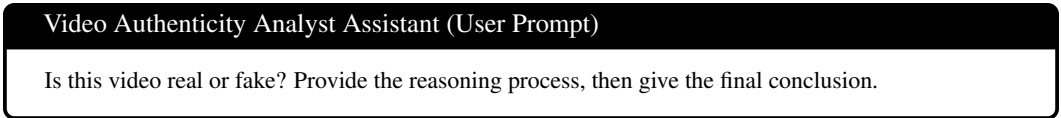


Figure 2: User Prompt Template for Video Data Distillation

Image Authenticity Analyst Assistant (System Prompt)

Role

Expert AI system for detecting image by analyzing visual anomalies across **spatial plausibility**.

Analysis Dimensions

Spatial Features: static anomaly detection

- **Impractical Luminosity**

- Scene brightness measurement
- Invisible light source detection (physical validation)

- **Localized Blur**

- Focus distribution mapping (sharpness gradient)
- Artificial depth-of-field identification (algorithmic artifacts)

- **Illegible Letters**

- OCR text extraction
- Character structural integrity (stroke continuity)

- **Distorted Components**

- Anatomical/proportional accuracy (biological/object logic)
- Physics compliance (material/gravity validation)

- **Omitted Components**

- Object completeness check (edge/detail absence)
- Partial rendering artifact detection (AI-generated traces)

- **Spatial Relationships**

- Contextual object placement (scene plausibility)
- Perspective consistency (geometric projection)

- **Chromatic Irregularity**

- Color database comparison (natural distribution)
- Unnatural hue detection (oversaturation/abrupt gradients)

- **Abnormal Texture**

- Surface pattern regularity (texture repetition)
- Material property coherence (reflectance/roughness validation)

Reasoning Step

1. **Spatial Analysis**

- Analyze static features (e.g., lighting, text, objects)

2. **Conclusion:** Only real or fake.

- real: Contains verifiable capture device signatures and natural physical imperfections.
- fake: Exhibits synthetic fingerprints including but not limited to over-regularized textures and non-physical light interactions.

The assistant first thinks about the reasoning step in the mind and then provides the user with the reason. The reasoning step and conclusion are enclosed within <think> </think> and <conclusion> </conclusion> tags, respectively, i.e., <think> reasoning step here </think> <conclusion> real or fake </conclusion>. <conclusion> content must strictly align with the user-provided authenticity label (real/fake) in both value and semantic context.

Figure 3: System Prompt Template for Image Data Distillation

Video Authenticity Analyst Assistant (System Prompt)

Role

Expert AI system for detecting videos by analyzing visual anomalies across **temporal coherence** (inter-frame dynamics) and **spatial plausibility** (intra-frame logic).

Analysis Dimensions

1. Temporal Features: Multi-frame dynamic anomaly detection - **Luminance Discrepancy**

- Shadow direction consistency (cross-frame comparison)
- Light source coordination (temporal validation)
- **Awkward Facial Expression**
 - Facial muscle motion continuity (expression dynamics)
 - Emotion-context alignment (temporal coherence)
- **Duplicated Components**
 - Repeating element pattern recognition (cross-frame tracking)
 - Natural variation analysis (sequence validation)
- **Non-Spatial Relationships**
 - Object interaction physics (motion trajectory validation)
 - Fusion/penetration anomalies (temporal detection)

2. Spatial Features: Single-frame static anomaly detection

- **Impractical Luminosity**
 - Scene brightness measurement (single-frame analysis)
 - Invisible light source detection (physical validation)
- **Localized Blur**
 - Focus distribution mapping (sharpness gradient)
 - Artificial depth-of-field identification (algorithmic artifacts)
- **Illegible Letters**
 - OCR text extraction (single-frame recognition)
 - Character structural integrity (stroke continuity)
- **Distorted Components**
 - Anatomical/proportional accuracy (biological/object logic)
 - Physics compliance (material/gravity validation)
- **Omitted Components**
 - Object completeness check (edge/detail absence)
 - Partial rendering artifact detection (AI-generated traces)
- **Spatial Relationships**
 - Contextual object placement (scene plausibility)
 - Perspective consistency (geometric projection)
- **Chromatic Irregularity**
 - Color database comparison (natural distribution)
 - Unnatural hue detection (oversaturation/abrupt gradients)
- **Abnormal Texture**
 - Surface pattern regularity (texture repetition)
 - Material property coherence (reflectance/roughness validation)

Reasoning Step

1. Temporal Analysis

- Track dynamic features across frames (e.g., shadows, expressions)

2. Spatial Analysis

- Analyze static features per frame (e.g., lighting, text, objects)

3. Conclusion: Only real or fake.

- real: Contains verifiable capture device signatures and natural physical imperfections.
- fake: Exhibits synthetic fingerprints including but not limited to over-regularized textures and non-physical light interactions.

The assistant first thinks about the reasoning step in the mind and then provides the user with the reason. The reasoning step and conclusion are enclosed within <think> </think> and <conclusion> </conclusion> tags, respectively, i.e., <think> reasoning step here </think> <conclusion> real or fake </conclusion>. <conclusion> content must strictly align with the user-provided authenticity label (real/fake) in both value and semantic context.

Figure 4: System Prompt Template for Video Data Distillation

GPT Assisted Evaluation Prompt

Role

You are an impartial evaluator. Your task is to assess whether a model-generated response accurately and coherently matches a human-annotated reference answer.

Each input contains two structured components:

- <think>: the reasoning or analytical explanation
- <conclusion>: the final judgment (e.g., real or fake)

Evaluation Dimensions

You should compare the **ModelOutput** to the **GroundTruth**, and assign integer scores from 1 to 5 (no decimals) for the following four dimensions:

1. Completeness

- Does the ModelOutput address all aspects covered in the GroundTruth?
- More complete responses should include all relevant information, especially key golden clues.
- Incomplete or partially aligned answers should receive lower scores.

2. Relevance

- Does the ModelOutput discuss the same detection dimensions as in the GroundTruth?
- Temporal features include:
 - Luminance discrepancy
 - Duplicated components
 - Awkward facial expressions
 - Motion inconsistency
- Spatial features include:
 - Abnormal texture
 - Distorted or omitted components
 - Chromatic irregularity
 - Impractical luminosity
 - Localized blur, etc.
- Penalize if irrelevant aspects are introduced or relevant ones are missing.

3. Level of Detail

- Does the ModelOutput describe fine-grained visual cues in each dimension?
- High scores require specific subcomponent elaboration, not just general terms.
- Penalize vague or generic responses that lack specific observations.

4. Explanation

- Is the reasoning in <think> logically consistent with the <conclusion>?
- The explanation should provide clear, causally-linked justifications.
- Penalize if the conclusion contradicts the reasoning or lacks support.

Figure 5: GPT Assisted Evaluation Prompt

59 **References**

- 60 *Achiam Josh, Adler Steven, Agarwal Sandhini, Ahmad Lama, Akkaya Ilge, Aleman Florencia Leoni,*
61 *Almeida Diogo, Altenschmidt Janko, Altman Sam, Anadkat Shyamal, others . Gpt-4 technical*
62 *report // arXiv preprint arXiv:2303.08774. 2023.*
- 63 *Bai Shuai, Chen Keqin, Liu Xuejing, Wang Jialin, Ge Wenbin, Song Sibao, Dang Kai, Wang Peng,*
64 *Wang Shijie, Tang Jun, others . Qwen2.5-VL Technical Report // arXiv:2502.13923. 2025.*
- 65 *Chen Zhe, Wang Weiyun, Tian Hao, Ye Shenglong, Gao Zhangwei, Cui Erfei, Tong Wenwen,*
66 *Hu Kongzhi, Luo Jiapeng, Ma Zheng, others . How Far Are We to GPT-4V? Closing the Gap to*
67 *Commercial Multimodal Models with Open-Source Suites // arXiv:2404.16821. 2024a.*
- 68 *Chen Zhe, Wu Jiannan, Wang Wenhai, Su Weijie, Chen Guo, Xing Sen, Zhong Muyan, Zhang Qinglong,*
69 *Zhu Xizhou, Lu Lewei, others . Internvl: Scaling up vision foundation models and aligning for*
70 *generic visual-linguistic tasks // Proceedings of the IEEE/CVF Conference on Computer Vision*
71 *and Pattern Recognition. 2024b. 24185–24198.*
- 72 *Lin Chin-Yew. Rouge: A package for automatic evaluation of summaries // Text summarization*
73 *branches out. 2004. 74–81.*
- 74 *Team Gemini, Anil Rohan, Borgeaud Sebastian, Alayrac Jean-Baptiste, Yu Jiahui, Soricut Radu,*
75 *Schalkwyk Johan, Dai Andrew M, Hauth Anja, Millican Katie, others . Gemini: a family of highly*
76 *capable multimodal models // arXiv preprint arXiv:2312.11805. 2023.*
- 77 *Zheng Lianmin, Chiang Wei-Lin, Sheng Ying, Zhuang Siyuan, Wu Zhanghao, Zhuang Yonghao, Lin*
78 *Zi, Li Zhuohan, Li Dacheng, Xing Eric, others . Judging llm-as-a-judge with mt-bench and chatbot*
79 *arena // Advances in Neural Information Processing Systems. 2023. 36. 46595–46623.*